

# Lab 1: Data Manipulation, Random Number Generation

July 7, 2020

Today's agenda: Manipulating data objects; using the built-in functions, doing numerical calculations, and basic plots; reinforcing core probabilistic ideas.

0. Open a new R Markdown file; set the output to HTML mode and “Knit”. This should produce a web page with the knitting procedure executing your code blocks. You can edit this new file to produce your homework submission.

## Background

The exponential distribution is defined by its cumulative distribution function

$$F(x) = 1 - e^{-\lambda x}$$

The R function `rexp` generates random variables with an exponential distribution.

```
rexp(n=10, rate=5)
```

```
## [1] 0.38340462 0.03866304 0.24814625 0.44634820 0.20847509 0.33775319
## [7] 0.04744719 0.19942565 0.05813347 0.14102286
```

produces 10 exponentially-distributed numbers with rate ( $\lambda$ ) of 5. If the second argument is omitted, the default rate is 1; this is the **standard exponential distribution**.

## Part I

1. Generate 200 random values from the standard exponential distribution and store them in a vector `exp.draws.1`. Find the mean and standard deviation of `exp.draws.1`.

```
exp.draws.1 <- rexp(200)
mean(exp.draws.1)
```

```
## [1] 0.9239318
```

```
sd(exp.draws.1)
```

```
## [1] 0.8429488
```

2. Repeat, but change the rate to 0.1, 0.5, 5 and 10, storing the results in vectors called `exp.draws.0.1`, `exp.draws.0.5`, `exp.draws.5` and `exp.draws.10`.

```
exp.draws.0.1 <- rexp(200, rate = 0.1)
mean(exp.draws.0.1)
```

```
## [1] 10.71301
```

```
sd(exp.draws.0.1)
```

```
## [1] 10.48842
```

```
exp.draws.0.5 <- rexp(200, rate = 0.5)
mean(exp.draws.0.5)
```

```
## [1] 2.268291
```

```
sd(exp.draws.0.5)
```

```
## [1] 2.114471
```

```
exp.draws.5 <- rexp(200, rate = 5)
mean(exp.draws.5)
```

```
## [1] 0.1956412
```

```
sd(exp.draws.5)
```

```
## [1] 0.1952145
```

```
exp.draws.10 <- rexp(200, rate = 10)
mean(exp.draws.10)
```

```
## [1] 0.09755564
```

```
sd(exp.draws.10)
```


```
## [1] 0.09568435
```

3. The function `plot()` is the generic function in R for the visual display of data. `hist()` is a function that takes in and bins data as a side effect. To use this function, we must first specify what we'd like to plot.


- a. Use the `hist()` function to produce a histogram of your standard exponential distribution.

```
hist(exp.draws.1) # or use ggplot2
```

**Histogram of exp.draws.1**




```
tibble(expdraw = exp.draws.1) %>% ggplot(aes(x = exp.draws.1)) +  
  geom_histogram(bins = 15)
```

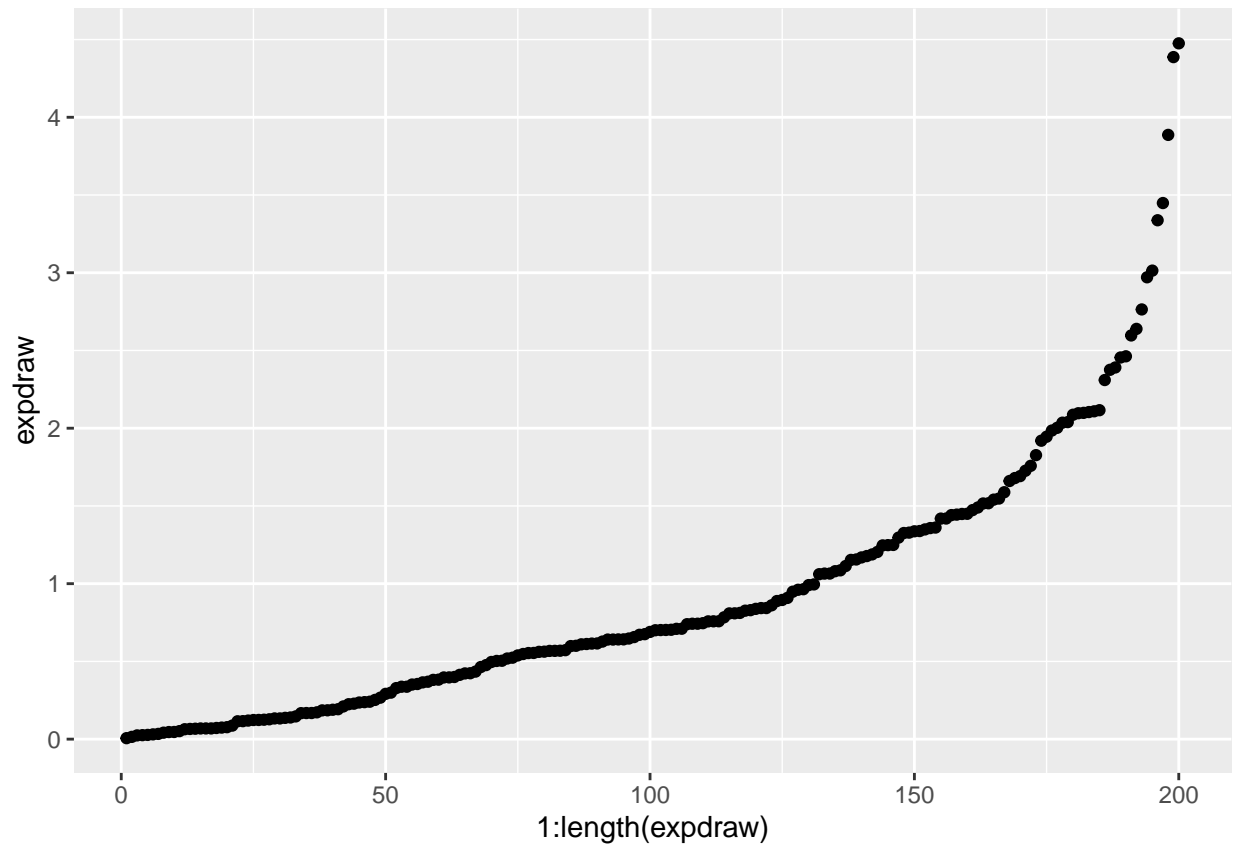


b. Use `plot()` with this vector to display the random values from your standard distribution in order.


```
plot(sort(exp.draws.1)) # or use ggplot2
```




```
tibble(expdraw = exp.draws.1) %>% arrange(expdraw) %>%  
  ggplot(aes(x = 1:length(expdraw), y = expdraw))+  
  geom_point()
```



c. Now, use `plot()` with two arguments -- any two of your other stored random value vectors -- to create a plot. For example, you could use `plot(exp.draws.1, exp.draws.5)` or use *ggplot2*.



```
tibble(x = exp.draws.1, y = exp.draws.5) %>%  
  ggplot(aes(x = x, y = y)) +  
  geom_point()
```



4. We'd now like to compare the properties of each of our vectors. Begin by creating a vector of the means of each of our five distributions in the order we created them and saving this to a variable name of your choice. Using this and other similar vectors, create the following scatterplots:
  - a. The five means versus the five rates used to generate the distribution.


```

exprn <- tibble(exp.1 = exp.draws.1, exp.0.1 = exp.draws.0.1, exp.0.5 = exp.draws.0.5, exp.5 = exp.draws.5, exp.10 = exp.draws.10)
emeans <- tibble(rates = c(1, 0.1, 0.5, 5, 10), means = colMeans(exprn), stds = apply(exprn, 2, sd))
rord <- order(emeans$rates)
emeans <- emeans[rord,]
# after chapter 4
emeans1 <- exprn %>% gather(key = 'draws', value = 'value') %>%
  group_by(draws) %>% summarize(means = mean(value), stds = sd(value)) %>% arrange(desc(means)) %>% ungroup()


## `summarise()` ungrouping output (override with `.groups` argument)
plot(means ~ rates, data = emeans)

```






```
# or using ggplot2  
emeans1 %>% ggplot(aes(x = rates, y = means))+  
  geom_point()
```




b. The standard deviations versus the rates.

```
plot(stds ~ rates, data = emeans)
```




```
# or using ggplot2  
emeans1 %>% ggplot(aes(x = rates, y = stds)) +  
  geom_point()
```




c. The means versus the standard deviations.

```
plot(means ~ stds, data = emeans)
```



```
# or using ggplot2  
emeans1 %>% ggplot(aes(x = stds, y = means))+  
  geom_point()
```



For each plot, explain in words what's going on.

## Part II

5. R's capacity for data and computation is large to what was available 10 years ago.

- a. To show this, generate 1.1 million numbers from the standard exponential distribution and store them in a vector called `big.exp.draws.1`. Calculate the mean and standard deviation.

```
big.exp.draws.1 <- rexp(1100000)
mean(big.exp.draws.1)
```

```
## [1] 1.000194
```


```
sd(big.exp.draws.1)
```

```
## [1] 0.9999471
```


- b. Plot a histogram of ``big.exp.draws.1``. Does it match the function  $1-e^{-x}$ ? Should it?

```
hist(big.exp.draws.1, probability = T)
afun <- function(x) 1-exp(-x)
curve(afun, add = T)
```

## Histogram of big.exp.draws.1



```
# or use ggplot2
tibble(expdraw = big.exp.draws.1) %>% ggplot()+
  geom_histogram(aes(x = big.exp.draws.1, y = (..count..)/sum(..count..)), binwidth = 0.5)+
  ylab('frequency')+
  stat_function(fun = afun, col = 'blue')
```



c. Find the mean of all of the entries in ``big.exp.draws.1`` which are strictly greater than 1. You may use

```
big2 <- big.exp.draws.1 > 1
mean(big.exp.draws.1[big2])
```

```
## [1] 2.000028
```

*# or use dplyr after chapter 4*

```
tibble(draw = big.exp.draws.1) %>% filter(draw > 1) %>% summarize(mean = mean(draw))
```

```
## # A tibble: 1 x 1
##   mean
##   <dbl>
## 1 2.00
```


d. Create a matrix, ``big.exp.draws.1.mat``, containing the the values in

`big.exp.draws.1`, with 1100 rows and 1000 columns. Use this matrix as the input to the `hist()` function and save the result to a variable of your choice. What happens to your data?

```
big.exp.draws.1.mat <- matrix(big.exp.draws.1, nrow = 1100)
bighist <- hist(big.exp.draws.1.mat)
```



## Histogram of big.exp.draws.1.mat



e. Calculate the mean of the 371st column of `big.exp.draws.1.mat`.

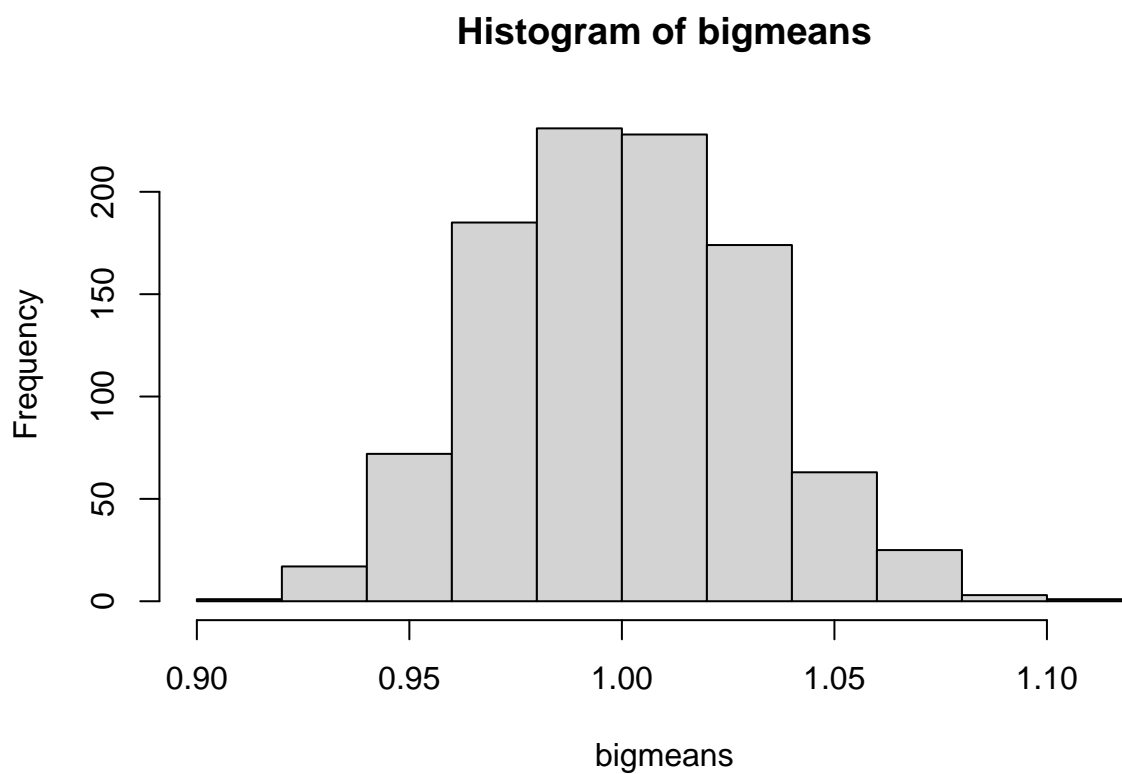
```
mean(big.exp.draws.1.mat[,371])
```

```
## [1] 1.011575
```

f. Now, find the means of all 1000 columns of `big.exp.draws.1.mat` simultaneously. Plot the histogram of

```
bigmeans <- colMeans(big.exp.draws.1.mat)
```

```
hist(bigmeans)
```



g. Take the square of each number in `big.exp.draws.1`, and find the mean of this new vector. Explain

```
bigsqrt <- sqrt(big.exp.draws.1)
mean(bigsqrt)
```

```
## [1] 0.8864046
```